

Building Containers for HPC

Charles Peterson



Overview

👋 Welcome!

In this workshop, we will explore the advanced use of containers on HPC resources, like UCLA's Hoffman2. This is a follow-up to a previous workshop “[Containers for HPC Resources](#)”

- Dive into advanced container topics
- Build containers tailored for HPC resources
- Got suggestions for upcoming workshops?
 - cpeterson@oarc.ucla.edu



Files for this Presentation

👀 Viewing the slides

https://github.com/ucla-oarc-hpc/WS_MakingContainers

Find this slide deck in the `slides` folder.

- html format: https://ucla-oarc-hpc.github.io/WS_MakingContainers/
- PDF format: `WS_MakingContainers.qmd`

Note

This presentation was created with [Quarto](#) and RStudio.

- Quarto file: `WS_MakingContainers.qmd`

Hands-on Exercise Setup

In this workshop, we will engage in hands-on exercises.

```
1 git clone https://github.com/ucla-oarc-hpc/WS_MakingContainers
```

Requirements:

- A computer with root/admin access (i.e., `sudo` permissions)
- Install [Apptainer](#)
- Install either [Docker](#) or [Podman](#)
 - In this workshop, we will be using Podman
- Note: Mac M1 (or other ARM-based systems) will not work for this exercise

Alternative: Virtual Machine

Pre-configured Virtual Machine (VM) available on [BOX](#)

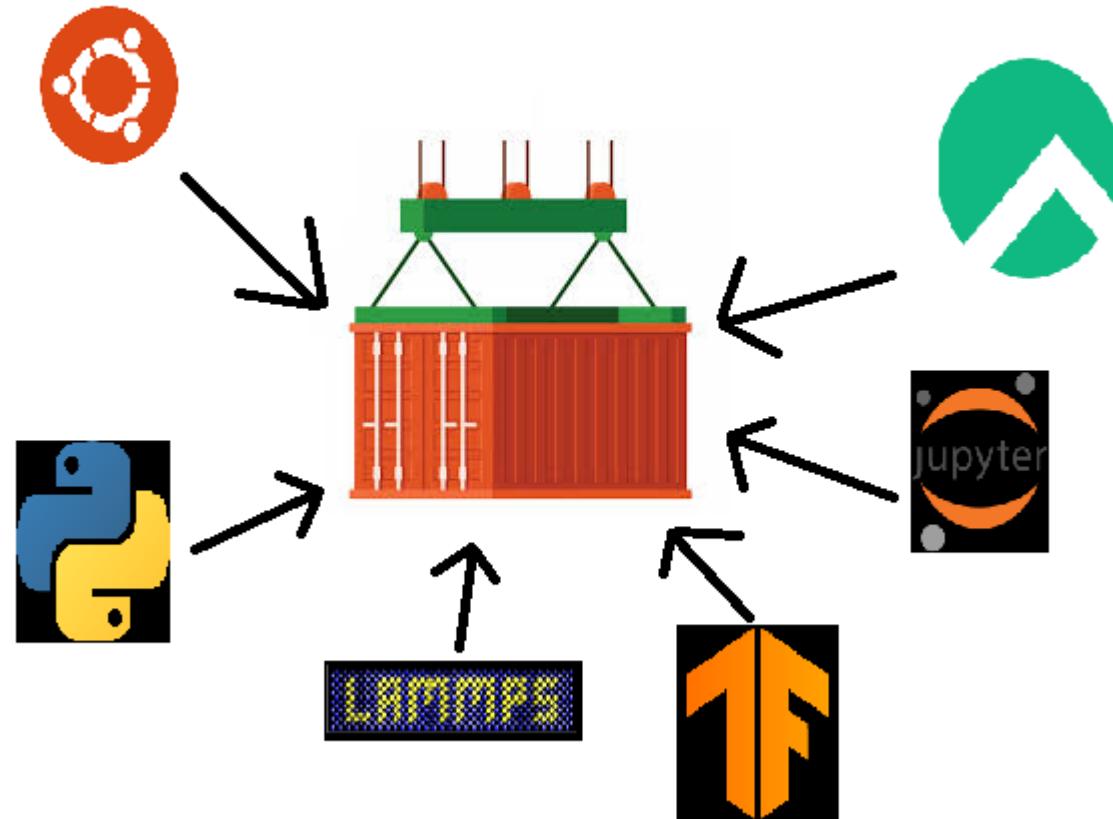
- VM file: [wscontainer.ova](#)
 - Recommended for use with VirtualBox
- Username/Password: wscontainer/wscontainer



Container Review



Containers



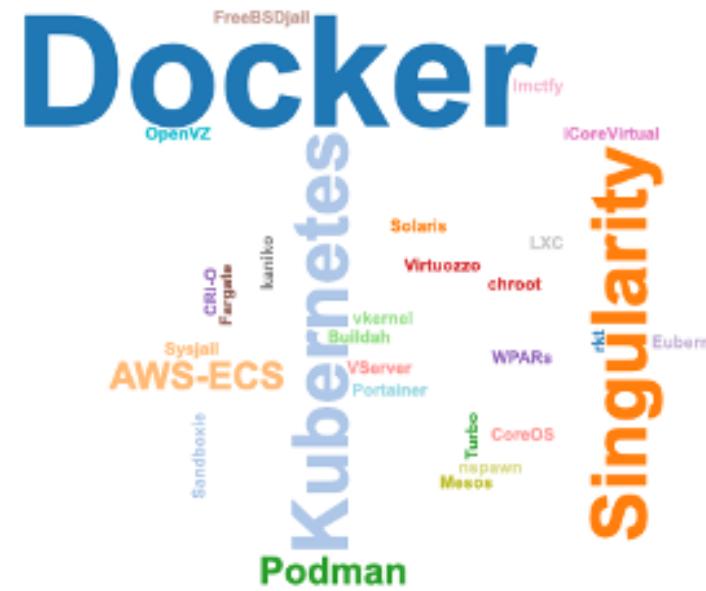
🛠️ Software for Containers

Apptainer

- Formerly Singularity
- 🎯 Designed and developed for HPC systems
- 🔐 Most likely installed on most HPC systems
- 💣 Supports Infiniband, GPUs, MPI, and other devices on the Host
- 🐳 Can run Docker containers



🛠 Software for Containers



Docker

- ⭐ Very popular
- ☁ Many popular cloud container registries
 - DockerHub, GitHub, Nvidia NGC
- 🚫 MPI not well supported
- 🛠 Most likely NOT available on many HPC systems



Software for Containers

Podman

- 🔒 Similar syntax as with Docker
 - 🔄 Can use to 'replace' Docker
- ❌ Doesn't have a root daemon process
- 🖥️ On some HPC resources (not on Hoffman2, yet)





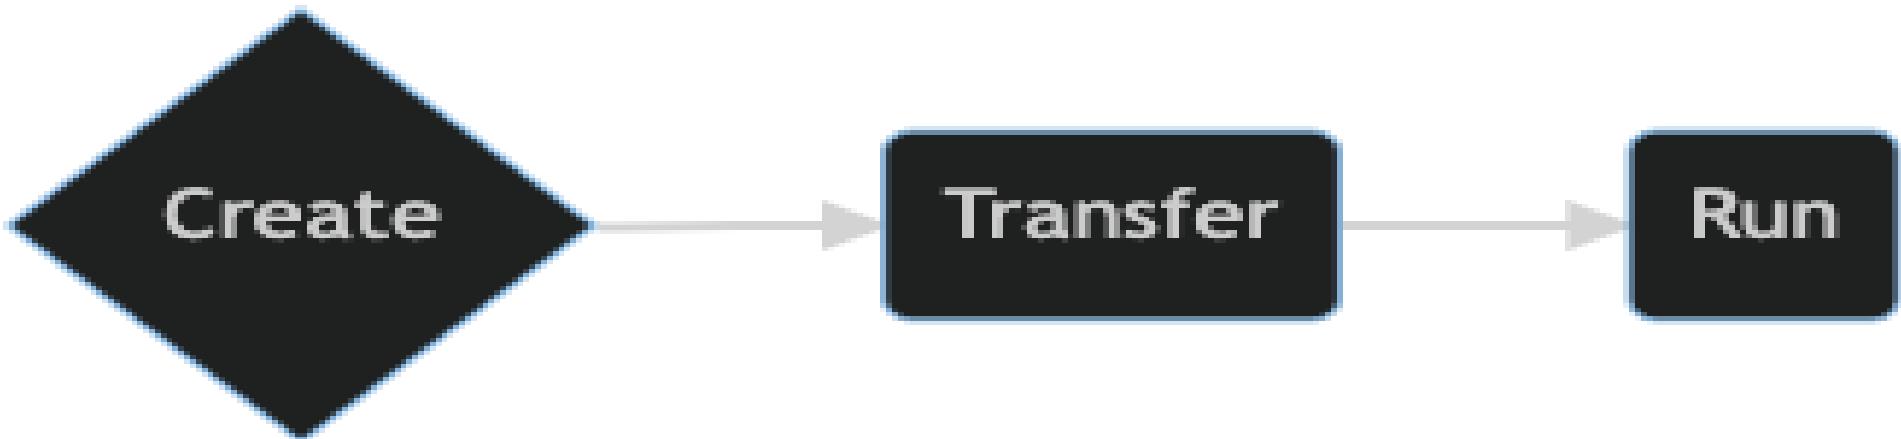
Apptainer workflow





Apptainer workflow (Create)

- 🏭 Build a container by installing Apptainer on your computer (where you have root/sudo access) to create a container
- 📬 Use a pre-built container
 - 🔎 Search Container Registries for container
 - DockerHub, GitHub packages, Nvidia NGC
- 💾 A SIF file contains the image for the container
- 📁 A sandbox container is a directory format used for writable containers





Apptainer workflow (Transfer)

🚚 Bring your container to Hoffman2

- 📈 Copy your container to Hoffman2

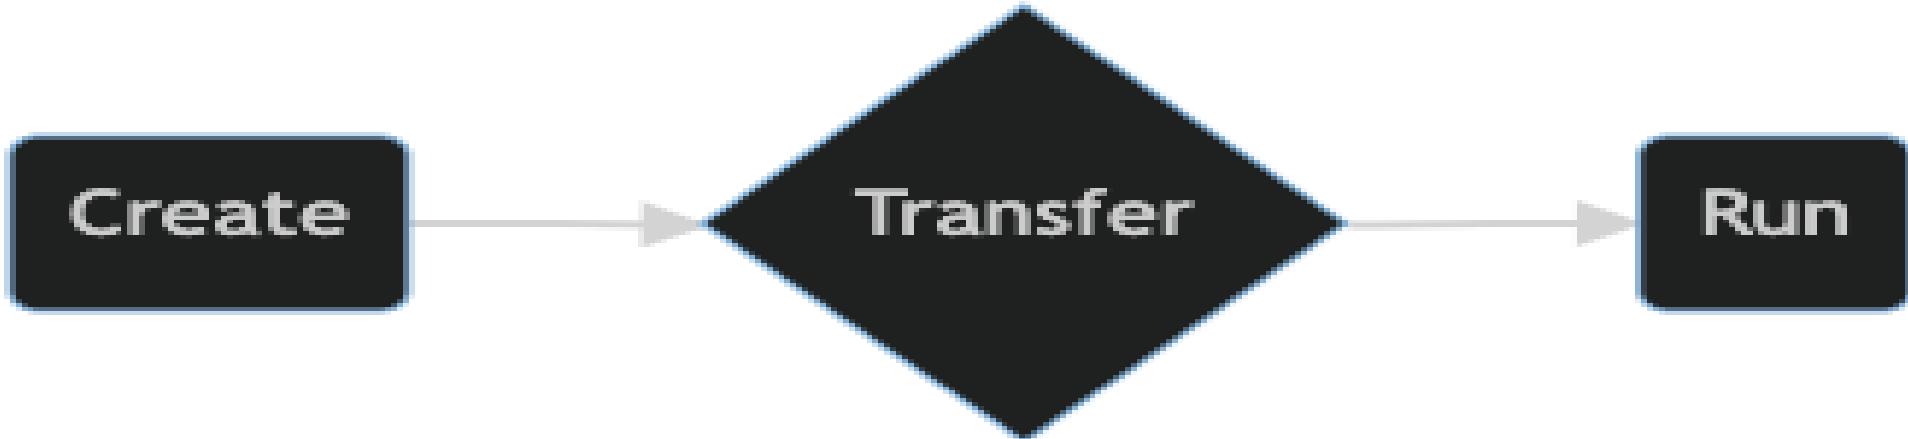
```
1 scp test.sif H2USERNAME@hoffman2.idre.ucla.edu
```

- 📈 Pull a container from online Container Register

```
1 apptainer pull docker://ubuntu:20.04
```

- 💻 Use a container pre-built on Hoffman2

```
1 module load apptainer
2 ls $H2_CONTAINER_LOC
```





Apptainer workflow (Run)

🤖 Interactive (qrsh) session

```
1 qrsh -l h_data=15G
2 module load apptainer
3 apptainer exec mypython.sif python3 test.py
```

📄 Batch (qsub) job

```
1 cat << EOF >> myjob.job
2 module load apptainer
3 apptainer exec mypython.sif python3 test.py
4 EOF
5 qsub -l h_data=15G myjob.job
```



⭐ Common Usage

On Hoffman2, to use apptainer, all you need to do is load the module

```
1 module load apptainer
```

- 🎉 Only module you need to load!
- No need to load tons of modules to run a single application
- 🌐 Expect MPI module if running parallel
 - `module load intel/2022.1.1`

📚 Common Apptainer commands

- 📤 Getting a container from somewhere
- 🏭 Build a container

```
1 apptainer pull [options]
2 apptainer pull docker://ubuntu:22.04
```

```
1 apptainer build [options]
2 apptainer build myapp.sif myapp.def
```

-  Run a single command inside of a container

```
1 apptainer exec [options]
2 apptainer exec myapp.sif MYCOMMAND
```

-  Run the the container with a predefined script

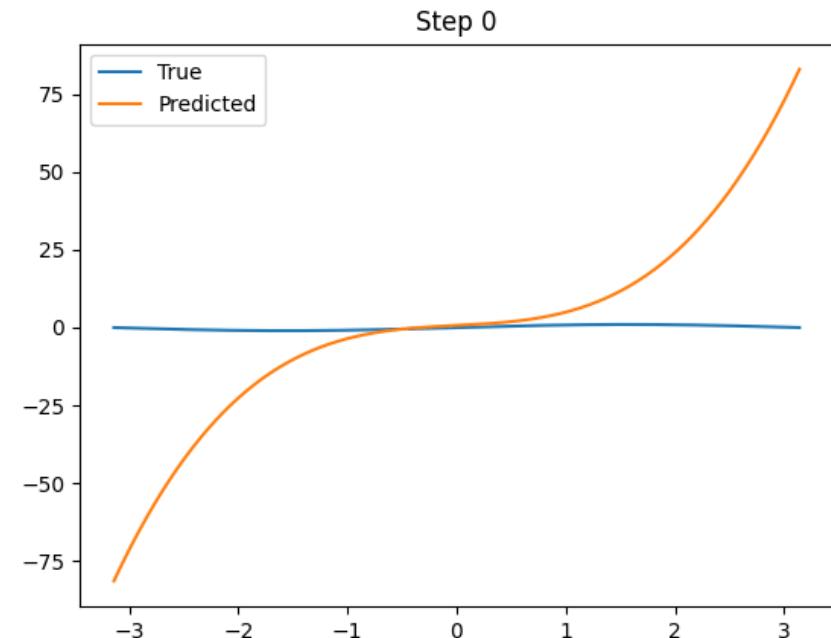
```
1 apptainer run [options]
2 apptainer run myapp.sif "SCRIPT ARGRMENTS"
```

Example 1: PyTorch

Example 1: Create writable containers



- This example uses PyTorch 🧠
- Similar to last week's example
 - Though, instead of using PyTorch's pre-build container, we will create our own container
- Go to the [EX1](#) directory
 - Examine the [pytorch.py](#) file
 - Optimize a 3rd order polynomial to a sine function



Creating the container 🚀

Sandbox image

- Create a container from a base source image
 - DockerHub has great minimal base containers
 - Ubuntu
 - Rocky Linux
 - CentOS
 - There are various versions of multiple builds
- Lets create a container based on Ubuntu 22.04

```
1 sudo apptainer build --sandbox ubuntu_22.04_SB/ docker://ubuntu:22.04
```

- sudo
 - Execute with root privileges
 - Building usually requires administrative access
- apptainer build:
 - Build new Apptainer container image
- --sandbox:
 - Flag for “sandbox” or a writable directory
- ubuntu_22.04_SB/
 - Name of the directory where the sandbox image will be created
- docker://ubuntu:22.04
 - Location of source image

Running the container



Next, we will start a WRITABLE interactive shell session in the sandbox image:

```
1 sudo apptainer shell --writable ubuntu_22.04_SB/
```

- **--writable** will allow us to modify the container

From here, we can run any commands we need to install PyTorch:

```
1 apt update
2 apt install -y python3 python3-pip
3 pip3 install torch torchvision torchaudio --extra-index-url https://download.pytorch.org/
4 exit
```

Convert the sandbox container to a SIF file, **pytorch.sif**

```
1 sudo apptainer build pytorch.sif ubuntu_22.04_SB/
```

Running Our Container



Run our SIF container:

```
1 apptainer exec pytorch.sif python3 pytorch.py
```

We do not need to be root since we are just running python3.

Transfer our container:

```
1 scp pytorch.sif hoffman2.idre.ucla.edu:
```

Example job script to run on Hoffman2:

```
1 qsub pytorch.job
```

In this example, we used an interactive approach to install a python package in a container.

This is a useful way to experiment installing of your applications

Example 2: Building from definition files

Creating containers

I coded a chemistry app located on github.

- Very early attempt at learning Hartree-Fock
 - Named from my hedgie, Quill Ricker
- <https://github.com/charliecpeterson/QUILL>

To install, we need

- Python with the PySCF package
- Eigen3, a Linear Algebra library



- Instead of installing these dependencies on H2 (or loding modules)
- Lets build a container!!

We will build this container by:

- Apptainter definition file (.def)
- Using Docker/Podman (Dockerfile)

Apptainer Definition file

Definition Files are like the blueprint to building a custom container.

Instead of interactively modifying a sandbox image, we can build a container with this Definition file

The [quill.def](#) file has all steps needed to build the QUILL container.

```
1 Bootstrap: docker
2 From: ubuntu:20.04
3
4 %labels
5 Author Charles Peterson <cpeterson@oarc.ucla.edu>
6
7 %post
8 apt-get update
9 DEBIAN_FRONTEND=noninteractive apt-get install -y --no-install-recommends
10    git python3 python3-dev python3-pip \
11    libeigen3-dev ca-certificates cmake make gcc g++
12 rm -rf /var/lib/apt/lists/*
13
14 pip3 install pyscf
15 ln -s /usr/bin/python3 /usr/bin/python
16 mkdir -pv /apps
17 cd /apps
18 git clone https://github.com/charliecpeterson/QUILL
19 cd QUILL
20 mkdir build ; cd build
21 cmake ..
22 make
23
24 %environment
25 export PATH=/apps/QUILL/build:$PATH
```

Sections:

- **Bootstrap/From** - Location of starting container
- **%labels** - adds metadata
- **%post** - This section runs commands to setup the final container
- **%environment** - define environment variables inside container

Create container

The `quill.sif` container is created

Let us test the container

- Run the command inside the container
 - `QUILL.x test.inp` inside the container
- Move container to Hoffman2

```
1 sudo apptainer build quill.sif quill.def
```

```
1 apptainer exec quill.sif QUILL.x test.inp
```

```
1 scp quill.sif H2USERNAME@hoffman2.idre.ucdavis.edu:
```

Building using Docker/Podman

- Docker or Podman can create containers
 - Then we convert Docker/Podman container to Apptainer
- I will be using Podman
 - Docker and Podman have same syntax
 - Replace **podman** with **docker**

```
1 docker build
2 docker images
3 docker pull
4 docker run
```

```
1 podman build
2 podman images
3 podman pull
4 podman run
```

Dockerfile

- Dockerfile-[quill](#) file is used by Docker to create container

```
1 FROM ubuntu:20.04
2
3 ## Author Charles Peterson <cpeterson@arc.ucla.edu>
4
5 RUN apt-get update \
6     && DEBIAN_FRONTEND=noninteractive apt-get install -y --no-install-recommends \
7     git python3 python3-dev python3-pip \
8     libeigen3-dev ca-certificates cmake make gcc g++ \
9     && rm -rf /var/lib/apt/lists/*
10
11 RUN pip3 install pyscf ; ln -s /usr/bin/python3 /usr/bin/python
12
13 RUN mkdir -pv /apps \
14     && cd /apps \
15     && git clone https://github.com/charliecpeterson/QUILL \
16     && cd QUILL \
17     && mkdir build ; cd build \
18     && cmake .. ; make
19
20 ENV PATH=/apps/QUILL/build:$PATH
```

Create container

Create Docker container

Build container from Dockerfile

See built docker (podman) container

Create Apptainer container

Save docker (podman) image in tarball

- Create SIF file
- Transfer to Hoffman2

```
1 podman build . -t quill:1.0 -f Dockerfile
```

```
1 podman image list
```

```
1 podman save quill:1.0 > quill.tar
```

```
1 apptainer build quill.sif docker-archive:
```

```
1 scp quill.sif H2USERNAME@hoffman2.idre.uc
```

Container Registry

In the previous slides, we created a SIF file (quill.sif), then transfer (scp) them to Hoffman2.

Instead of this, we can upload our container to a **Container Registry**.

- These **Registries** are used store our containers on a remote, cloud server that can then be pulled/download anywhere that has apptainer.
 - DockerHub
 - GitHub Packages

DockerHub

Lets create a repo on DockerHub

- First, create a DockerHub account
 - Mine is [charliecpeterson](#)
- Push our podman container to DockerHub
 - Registry location [docker.io](#)
- Tag container to DockerHub location
- Login info to DockerHub
- Push container to DockerHub
- Then pull the container on Hoffman2

```
1 podman tag quill:1.0 docker.io/charliecpeterson
```

```
1 podman login docker.io
```

```
1 podman push docker.io/charliecpeterson/quill:1.
```

```
1 apptainer pull docker://docker.io/charliecpeter
```

GitHub Packages

- Lets create a repo on GitHub
 - Look for the Packages tab
- Same syntax as before
 - registry location is `ghcr.io`

Push our final container to GitHub

```
1 podman tag quill:1.0 ghcr.io/charliecpeterson/quill:1.0
2 podman push ghcr.io/charliecpeterson/quill:1.0
```

Then pull the container on Hoffman2

```
1 apptainer pull docker://ghcr.io/charliecpeterson/quill:1.0
```

DockerHub and GitHub Packages are popular cloud registries. You can create and deploy a local container registry.

- <https://docs.docker.com/registry/deploying/>

Running Container

Once the container is on Hoffman2, submit job.

```
1 #!/bin/bash
2 #$ -cwd
3 #$ -o quill.$JOB_ID
4 #$ -j y
5 #$ -l h_rt=1:00:00,h_data=15G
6 #$ -pe shared 1
7 #$ -l arch=intel-gold*
8
9 # load the job environment:
10 . /u/local/Modules/default/init/modules.sh
11 module load apptainer
12
13 # Container part: apptainer exec QUILL.sif
14 # Command: QUILL.x /apps/QUILL/input.inp
15 time apptainer exec quill.sif QUILL.x test.inp
```

Submit job script

```
1 qsub test.job
```

More information on using Definition files

- https://apptainer.org/docs/user/1.0/definition_files.html

More information on using Dockerfiles

- <https://docs.docker.com/engine/reference/builder/>

Example 3: Anaconda

Anaconda is a very popular python and R distribution for simplifying package installation

- Check out my [Hoffman2 Happy Hour on Anaconda](#) on our GitHub workshop page

Though Anaconda can be tricky installing in a container due to environment setup.

We will have an example using Anaconda to install an application in a container.

Building H2O

We will go over creating a definition file for a example with Anaconda.

We will install the software [h2o.ai](https://www.h2o.ai). This is a great machine learning platform that has Python and R libraries.

In this example, we will use Anaconda to install h2o packages inside python and R.

H2o definition file

The `h2o.def` file

- Definition file for Apptainer
- Install Anaconda from an `env.yml`
- Use of `%runscript` to setup Anaconda env for `apptainer run`
 - the `$@` take arguments as a string from the command line
 - `apptainer run h2o.sif "python h2o-test.py"`

```
1 Bootstrap: docker
2 From: ubuntu:22.04
3
4
5 %labels
6 Author Charles Peterson <cpeterson@oarc.ucla.edu>
7
8
9 %post
10 export DEBIAN_FRONTEND=noninteractive
11 apt -y update ; apt -y upgrade
12 apt install -y wget libbz2-dev wget git gcc libreadline-dev zlib1g-dev default-jre default-jdk
13
14 #Install anaconda
15 cd /tmp
16 wget https://repo.anaconda.com/archive/Anaconda3-2022.05-Linux-x86_64.sh
```

```
17 bash Anaconda3-2022.05-Linux-x86_64.sh -b -p /opt/anaconda
18 bash -c "source /opt/anaconda/etc/profile.d/conda.sh
19 conda create -n h2oai h2o -c h2oai -c conda-forge
20 "
21
22 %runscript
23 exec bash -c "source /opt/anaconda/etc/profile.d/conda.sh
24 conda activate h2oai
25 $@"
```

Building container

- Create h2o.sif
- Run h2o.R inside the container

```
1 sudo apptainer build h2o.sif h2o.def
```

```
1 apptainer run h2o.sif "python h2o-test.py"
```



Note

- **apptainer exec foo.sif [COMMAND]**
 - Run a single [COMMAND] inside the container
 - The runscript will NOT run
- **apptainer run foo.sif**
 - Run the runscript inside the container

Example 4: RStudio and Jupyter

Jupyter

- Create a container with Jupyter
 - Useful so you can run same Jupyter setup anywhere
 - Install all packages you need in container
- `jupyter.def`
 - Apptainer Def file
 - Starts with a pre-built python 3.8.13 container
 - Adds jupyter, pandas, and seaborn
- Build container (locally)
- Transfer to H2
- Start Jupyter on Hoffman2
 - Note the name of compute node you landed on

```
1 sudo apptainer build jupyter.sif jupyter.def
2 scp jupyter.sif hoffman2.idre.ucla.edu:
```

```
1 qrsh -l h_data=10G
2 module load apptainer
3 hostname
4 apptainer exec jupyter.sif jupyter lab --ip 0.0.0.0
```

- SSH tunnel to H2 compute node

- Change nXXX (compute node)
- Change port 8888 if needed

```
1 ssh -L 8888:nXXX:8888 username@hoffman2.idre.ucla.edu
```

- Then open a web browser and type

```
1 http://localhost:8888
```

Rstudio

- Use Rstudio Server to open a Rstudio session on your web browser.
 - Very similar to the Jupyter setup
- More information from a previous [H2HH Rstudio session](#)
- Create this container using Docker/Podman
 - Local computer
- Setup Rstudio on Hoffman2
- Start Rstudio on web browser

```
1 podman build -f Dockerfile-rstudio -t rstudio:4.1.0 .
2 podman save rstudio:4.1.0 > rstudio.tar
3 apptainer build rstudio.sif docker-archive://rstudio.ta
4 scp rstudio.sif hoffman2.idre.ucla.edu
```

```
1 # get an interactive job
2 qrsh -l h_data=10G
3 # Create tmp directories
4 mkdir -pv $SCRATCH/rstudiotmp/var/lib
5 mkdir -pv $SCRATCH/rstudiotmp/var/run
6 mkdir -pv $SCRATCH/rstudiotmp/tmp
7 #Setup apptainer
8 module load apptainer
9 #Run rstudio
10 apptainer run -B $SCRATCH/rstudiotmp/var/lib:/var/lib/
11 # This command will display some information and a `ssl`
```

```
1 ssh -L 8787:nXXX:8787 username@hoffman2.idre.ucla.edu
2 # Then open a web browser and type
3 http://localhost:8787 #or whatever port number that was
```

Things to Think About



Tips



Size of container

- Try to keep the size of your container small and minimal
 - Only have the things necessary for your applications to run
- Large containers will need more **memory** and will take more time to start up
- Check out [Multi-Stage](#) approach
- Build from an existing container
 - Look for images on DockerHub, NGC
 - Build your own and upload to DockerHub/GitHub
- Good idea to build a sandbox container, then create definition file
 - Test out commands in sandbox while making the def file!

More Things to Think About



- Share .sif files with your friends!
 - Save your (Docker) containers to DockerHub or GitHub Packages
 - Create Dockerfile/Def files to recreate or modify containers
- Find examples of Dockerfiles and Apptainer def files on our GitHub Page
 - https://github.com/ucla-oarc-hpc/hpc_containers
- Use DMTCP Checkpointing with Containers!!
 - <https://apptainer.org/docs/user/main/checkpoint.html>

Thank you! 🙏

Questions? Comments? 🤔

Charles Peterson cpeterson@oarc.ucla.edu



